

# Spatio-Temporal Site Recommendation

Blinded for Double-Blind review

**Abstract**—Recommendation systems have become extremely common in recent years, and are utilized in a variety of areas to predict the “rating” or “preference” that a user would give to a point of interest (PoI), such as a restaurant, a hotel, or a bar. Such systems typically produce a list of recommendations by considering previous ratings of the user, as well as ratings of other users. Not every person rates every point of interest they visit. In this work, we want to explore the use of spatio-temporal data to improve recommendation systems: We postulate that spatio-temporal user data may indicate the liking or disliking of a point of interest. Clearly, if a user frequently visits the same PoI, stays at the PoI for long times, and is willing to travel a long distance to visit a PoI, that might indicate that user likes that PoI. Thus, we propose to extract user-PoI relation features from spatio-temporal trajectory data only. Using these features, we use out-of-the-box data mining and machine learning solutions, to estimate the popularity of a PoI. Our experimental evaluation shows, that the features extracted from spatio-temporal data able to accurately predict the popularity of a PoI, using ground-truth data from Yelp as a baseline.

## I. INTRODUCTION

Modern technology to capture geo-spatial information produces a huge flood of individual trajectory data, coupled with a new user mentality of utilizing this technology to voluntarily share information. By mining this data, and thus turning it into actionable information, the McKinsey Global Institute [5] projects a “\$600 billion potential annual consumer surplus from using personal location data globally”. Towards this goal of making location data actionable, we propose to mine potential user ratings for location sites (e.g., restaurants, shops) from spatial-temporal data. User rating is critical in many applications, especially recommendation system. Techniques such as widely used collaborative filtering use known user rating information to estimate the interests of a user. However, in practice, the user-site rating matrix is usually highly sparse, meaning that there are not enough user rating information to perform such tasks. This is known as the cold-start problem. There are many reasons, for example, many users do not want to spend time to rate locations they visited, or locations such as shops do not have an efficient way to ask users for feedback.

To solve this problem, we propose to obtain implicit user-rating rather than explicit user ratings. Our approach uses trajectory data to find users that have likely visited a site. Since this data does not explicitly tell us whether the user likes that site, we propose to mine features from the users trajectory, which we intuitively expect to implicitly describe

whether the user likes that site. The features that we propose to obtain from trajectory data include:

- **The frequency of visits of a user.** If a user visits a site only once in their life, chances are that the user did not like the site enough to return. If the site is frequented often by the user, he seems to like it.
- **The length of stay of a user.** If a user stays at a restaurant or a hotel for a long time, that indicates that he likes the site.
- **The distance to their home base.** If a user is willing to take a long journey to reach a site, thus bypassing other, similar sites, that indicates that the user is subject to a strong attraction from that site, indicating that the user likes that site.

There are several advantages of using location data for location rating: (i) A potential solution to the cold-start problem, (ii) no user effort to capture their recommendation such as filling in rating forms, and (iii) it is based on user’s behavior, thus more objective and prone to alteration by fake-user-ratings and spam/bot-user-ratings.

Our approach becomes viable, due to the abundance of large open-sources collections of voluntarily contributed trajectory data, including the following data sources:

- **Location-Based Social Networks (LBSNs)** allows user to “Check-in” into a physical site such as a hotel, a restaurant or a metro station. Fairly large LBSN datasets, made anonymous, are available publicly. For instance, the FourSquare dataset used in [8] contains more than 30 million checkins and is available publicly. Such data explicitly includes the sites that a user has visited.
- **Geocoded Social Media Data:** is obtainable from public streaming APIs including for Twitter, Instagram, and Flickr. These data sources provide low-frequency trajectory data. Yet, microblogs and images are often published in sites of interest to a user, thus giving implicit information about the users site preferences.
- As part of the effort to create **Open-Street-Map (OSM)** [6], [2] road network data, users have been uploading GPS traces of their routes to the OSM site. While these routes are typically used to digitize the road network, the majority of routes is uploaded by pedestrians, and can be used as an indicator of sites that the corresponding user frequents. These trajectories are publicly available through the OSM API.

To describe our approach of site-recommendation using trajectory data, the rest of this work is organized as follows. We survey the state-of-the-art on site-recommendation and location-based recommendation systems in Section II. Then, we formally define the problem location-based site recommendation in Section III. Our solution, using deep-learning to bridge the gap from user-behaviour to user-site-recommendations is given in Section V. Our solution is evaluated in Section VII, showing that our rating prediction for a restaurant is able to closely predict authoritative ground-truth site-ratings obtained from Yelp. We conclude our work in Section VIII.

## II. RELATED WORK

### III. PROBLEM DEFINITION

In this section we formally define a user-trajectory, and define our notion of a user-site-stay, which we are going to use to extract features to estimate the affinity between a user and a site later in Section V. We first start by defining a trajectory as follows.

*Definition 1 (Spatio-Temporal Database):* Let  $\mathcal{U}$  denote a set of unique user identifiers, let  $\mathcal{G} = [-90, 90] \times [-180, 180]$  denote the space of longitude/latitude geo-coordinates, and the  $\mathcal{T}$  denote the time domain. A *spatio-temporal database*  $\mathcal{ST} \subseteq \mathcal{U} \times \mathcal{G} \times \mathcal{T}$  is a collection of triples  $(id \in \mathcal{U}, (lat, long) \in \mathcal{G}, t \in \mathcal{T})$ . Each triple  $(u, s, t) \in \mathcal{ST}$  is called an observation.

We group a spatio-temporal database into observations of the same user, denoted as user-trajectory, formally:

*Definition 2 (User-Trajectory):* Let  $\mathcal{ST}$  be a spatio-temporal database and let  $u \in \mathcal{U}$  be a user. The set

$$\mathcal{ST}(u) := \{(u', (lat, long), t) \in \mathcal{ST} | u = u'\}$$

is called the user-trajectory of user  $u$ .

In order to obtain recommendation information from a user-trajectory, we need to link the user-trajectory to sites, such as restaurants and hotels. For this purpose, we join a spatio-temporal database with a database of points of interest (such as provided by Open-Street Map) like restaurants and hotels. Next, we define our notion of a *stay*. A stay is an event of user visiting a site, enriched by the duration of the stay.

*Definition 3 (Stay Trajectory):* Let  $\mathcal{ST}$  be a spatio-temporal database and let  $\mathcal{S} \subseteq \mathcal{G}$  be a collection of  $(lat, long)$  pairs of sites. A stay is a triple  $(u \in \mathcal{U}, s \in \mathcal{S}, (t_{start}, t_{end}) \in \mathcal{T} \times \mathcal{T})$ , indicating that user  $u$  has stayed at site  $s$  from time  $t_{start}$  to time  $t_{end}$ . We let  $(\mathcal{ST} \bowtie \mathcal{S})$  denote the set of all stays mined from all trajectories  $\mathcal{ST}$  using all sites in  $\mathcal{S}$ . The sequence of all stays a user  $u \in \mathcal{U}$  is called the stay-trajectory  $(\mathcal{ST} \bowtie \mathcal{S})(u)$  of  $u$ , defined as:

$$(\mathcal{ST} \bowtie \mathcal{S})(u) := \{(u, p, (t_{start}, t_{end})) \in \mathcal{S} | u = u'\}$$

Finding stay points in trajectory and PoI databases is a research topic that has raised attention in the past. For instance, a state-of-the-art approach [4], [11], [10], [7] uses a distance threshold  $\theta_d$  and defines a stay as the duration of time where the trajectory does not exceed a distance of  $\theta_d$  to a PoI. In this work, we assume that *stay* detection algorithms are already

applied to a trajectory, such that user-trajectory is mapped to a sequence of PoI stays. This assumption is discussed in Section ??.

Given stay-trajectories for each user, the challenge of this work is to predict the rating between users and a site. As site is PoI which can be rated by a user, such as a restaurant or a hotel. We assume that we have a recommendation database, where users can rate sites. We assume normalized rating values in the interval  $[0, 1]$ , where 0 corresponds to the lowest rating and 1 corresponds to the highest rating.

*Definition 4 (User-Site Recommendation Database):* A recommendation database  $\mathcal{R}$  is a set of user ratings  $\mathcal{R} \subseteq \mathcal{U} \times \mathcal{S} \times [0, 1]$ .

The task of this work is to predict the triples in  $\mathcal{R}$ . That is, the challenge is to predict the rating that a user  $u \in \mathcal{U}$  will give to a site  $s \in \mathcal{S}$ , using a spatio-temporal given in  $\mathcal{ST}$ .

## IV. DISCUSSION: STAY POINT DETECTION

### V. SPATIO-TEMPORAL USER-SITE FEATURE EXTRACTION

There are many spatio-temporal features that could potentially be related to a user's feeling of a location. We discuss in the section a set of common features we select for the rating prediction problem, and the methods we use to extra such features. It is easy to understand why these features are selected since they are intuitive.

#### A. Extracting the frequency of visits of a user

Our assumption is, if a user repeatedly visits a PoI, e.g., always dine in the same restaurant, it strongly suggests that the user favours the PoI. On the other hand, if a user visits a PoI only once and never comes back, it suggests the user dislike the place. Therefore we choose the frequency of visits as a feature.

In order to extra frequency information from a user's trajectories, we count for how many time does the user visit a location in a given time period. Typically, each check-in to the location is counted as one visit. Users are then assigned into several frequency groups (Table I) based on how often they visit the location.

TABLE I: Groups based on how often user visits a location

Frequency groups
At least one visit per day
At least one visit per week
At least one visit per month
Less than one visit per month
Never visited

We note a user's behaviour may vary over time. For example, if a user likes a shopping center, he may visit the shopping center very frequently during Thanksgiving and similar holidays, while visit the same place only once in several month during the rest of the year. As a result, the timing window used to count visits can have significant impact on the user's group assignment. For the same user, if we consider only visits happened in the Thanksgiving week, the

user belongs to the “At least one visit per week ” group. However, if we look at the year-long visits, he may be grouped into “Less than one visit per month” since his visits are averaged over twelve months. We propose a simple way to mitigate this problem.

First, we use timing windows with different sizes simultaneously to compute the frequency of visits of a user. Specifically, given a user’s trajectory data for a period of  $n$  days, for each day/week/month within these  $n$  days, we count the number of visits that falls in the same day/week/month. We then use this count to compute the user’s group of that specific day/week/month. Here we use day/week/month as nature timing windows, but in practice the size of timing window can be arbitrary. Second, when a user’s visit frequency demonstrates inconsistency in different time period with the same timing window size, we use his maximal visiting frequency, e.g., if a user visits a location 3 times a week in week 1, but 0 times in week 2, the user is then categorized as “At least one visit per week” in this case.

Additionally, for each user, we calculate three numerical features that describe the user’s visiting pattern to a location: 1) **Average duration between two consecutive visits**, 2) **Minimal duration between two consecutive visits**, and 3) **Maximal duration between two consecutive visits**. These information are supplemental to the frequency groups. We are interested in finding out their impact on the rating prediction result.

#### B. Extracting the length of stays of a user

#### C. Extracting the travel distance of a user from its home base

If a user is willing to travel a long distance from his home/work area to a PoI, it is likely that the PoI is attractive to him. Unfortunately, measuring the travel distance from one’s home/work can be challenging. This is because users’ home or work address or coordinates are usually not explicitly marked in a spatial-temporal dataset due to privacy concerns.

We use distance-based location clustering (e.g., weighted k-means [3], where the weight of a location is the number of visits by the user) to estimate a user’s home base, i.e., an area where he is likely to live/work. Locations a user visits frequently is likely to formulate a dense cluster in the area where he lives and works [1], [9]. If there exists a location that is isolated from any of these clusters, we deem such a location as “far from home”. It requires extra effort for the user to travel this location. Willing to make such effort indicates the user may like the place. An example of home base and isolated locations is illustrated in Figure 1.

For each isolated location, we estimate the travelling effort for the user to reach the location. Using the distance between the isolated location and the user’s home base may be biased. A user who likes to drive can easily travel to a mall several miles away from home, but for someone who does not have a car, travelling the same distance means significantly more effort. Instead, we use the relative travelling distance. For a location  $l$ , we compute  $D_l = d_l / \overline{d_{home}}$  where  $d_l$  is the minimal travel distance between  $l$  and any home base location,

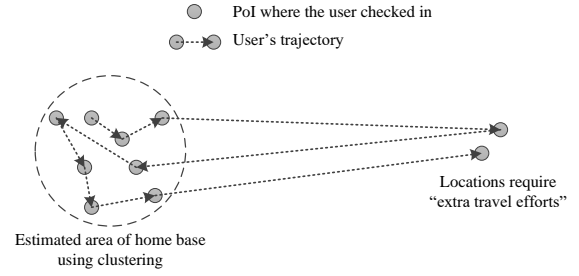


Fig. 1: Example of home base and isolated locations

and  $\overline{d_{home}}$  the average travel distance between locations within the user’s home base.

#### D. Extracting the total number of different visitors

In reality, a highly-rated location usually attracts many different users.

#### E. Extracting the type of location

#### F. Discussion

1. Assumptions: when are valid and when is not
2. Lack of data:
3. Explore social connection
4. Compare nearby PoIs

### VI. USER-SITE RATING PREDICTION

1. Simple linear regression - not enough expression power
2. Cov-DNN - Why?

### VII. EXPERIMENTAL EVALUATION

Schemes 1. Linear regression 2. Matrix Factorization - parameters? - Microsoft? 3. Cov-DNN - parameters?

Metrics 1. MSE of prediction 2. Running time(?) 3.

Dataset 1. NYC 2. Tokyo

Ground truth 1. FourSquare 2. Yelp 3. TripAdvisor

Parameters 1. Training set size / percentage 2. Category of locations 3. NN depth

Data cleaning 1. Being rated by more than 20 users 2.

### VIII. CONCLUSIONS

#### REFERENCES

- [1] P. Golle and K. Partridge. On the anonymity of home/work location pairs. In *International Conference on Pervasive Computing*, pages 390–397. Springer, 2009.
- [2] M. Haklay and P. Weber. Openstreetmap: User-generated street maps. *IEEE Pervasive Computing*, October-December:12–18, 2008.
- [3] K. Kerdprasop, N. Kerdprasop, and P. Sattayatham. Weighted k-means for density-biased clustering. In *International Conference on Data Warehousing and Knowledge Discovery*, pages 488–497. Springer, 2005.
- [4] Q. Li, Y. Zheng, X. Xie, Y. Chen, W. Liu, and W.-Y. Ma. Mining user similarity based on location history. In *Proceedings of the 16th ACM SIGSPATIAL international conference on Advances in geographic information systems*, page 34. ACM, 2008.
- [5] J. Manyika, M. Chui, B. Brown, J. Bughin, R. Dobbs, C. Roxburgh, and A. H. Byers. Big data: The next frontier for innovation, competition, and productivity. 2011.
- [6] Open street map. <http://www.openstreetmap.org>.
- [7] X. Xiao, Y. Zheng, Q. Luo, and X. Xie. Finding similar users using category-based location history. In *Proceedings of the 18th SIGSPATIAL International Conference on Advances in Geographic Information Systems*, pages 442–445. ACM, 2010.

- [8] D. Yang, D. Zhang, L. Chen, and B. Qu. Nationtelescope: Monitoring and visualizing large-scale collective behavior in lbsns. *Journal of Network and Computer Applications*, 55:170–180, 2015.
- [9] H. Zang and J. Bolot. Anonymization of location data does not work: A large-scale measurement study. In *Proceedings of the 17th annual international conference on Mobile computing and networking*, pages 145–156. ACM, 2011.
- [10] Y. Zheng, X. Xie, and W.-Y. Ma. Geolife: A collaborative social networking service among user, location and trajectory. *IEEE Data Eng. Bull.*, 33(2):32–39, 2010.
- [11] Y. Zheng, L. Zhang, X. Xie, and W.-Y. Ma. Mining interesting locations and travel sequences from gps trajectories. In *Proceedings of the 18th international conference on World wide web*, pages 791–800. ACM, 2009.